

RLMstep

Saar Yalov

December 26, 2017

The RLMstep algorithm is based on Robust Stepwise Regression (2000) by Claudio Agostinelli

1 Comparison of the RLMstep and stepAIC

The two methods yielded in completely different models. The RLMstep method gave the following model:

```
> RLMstep3(HtVol$HtVol,model.matrix(~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)
,data=HtVol),f.to.enter = 2,f.to.leave = 1,is.bca = T)
```

\$coef

data.in.model	data.in.modelHt:BSA	data.in.modelMale
36.031775	2.086633	30.793939

However stepAIC yields a much different model:

```
> stepAIC(lm(HtVol~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2),data=HtVol),direction
= "both")
```

Call:

```
lm(formula = HtVol ~ Age + Ht + Wt + BMI + BSA + Male + I(Age^2) +
I(Wt^2) + I(BMI^2) + I(BSA^2) + Age:Ht + Age:BMI + Age:BSA +
Age:Male + Ht:Wt + Ht:BMI + Ht:BSA + Ht:Male + Wt:BMI + Wt:BSA +
Wt:Male + BMI:Male + BSA:Male, data = HtVol)
```

Coefficients:

(Intercept)	Age	Ht	Wt	BMI
2.343e+04	8.179e+01	-3.239e+02	3.058e+02	-1.327e+03
BSA	Male	I(Age^2)	I(Wt^2)	I(BMI^2)
-9.345e+03	2.650e+03	-1.926e-02	-4.107e+00	1.357e+01
I(BSA^2)	Age:Ht	Age:BMI	Age:BSA	Age:Male
-4.798e+04	-8.639e-01	-1.762e+00	6.587e+01	-4.233e+00
Ht:Wt	Ht:BMI	Ht:BSA	Ht:Male	Wt:BMI
-8.669e+00	1.259e+01	5.240e+02	-3.590e+01	-1.011e+01
Wt:BSA	Wt:Male	BMI:Male	BSA:Male	
1.065e+03	-2.328e+01	-7.322e+01	4.323e+03	

To compare the fitness of the two models we will look at a scatter plot of their fitted values vs true value

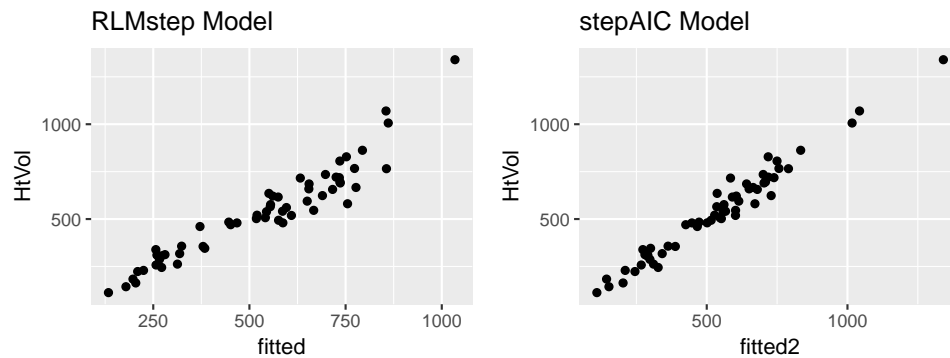
#RLM

```
> fitted1 <- rlm1$fitted.values
> p1<- ggplot( data= HtVol,aes(x=fitted1, y=HtVol))+geom_point()+
  ggtitle("RLMstep Model")
> dev.print(device=pdf, file="Stp598RLMmodel.pdf")
RStudioGD
2
```

#AIC

```
> fitted2 <- rlm2$fitted.values
> p2<- ggplot( data= HtVol,aes(x=fitted2, y=HtVol))+geom_point()+
  ggtitle("stepAIC Model")
> print(p2)
> dev.print(device=pdf, file="Stp598AICmodel.pdf")
RStudioGD
2
```

Figure 1: Residual Box Plots



The stepAIC seems to have more accurate predictions in comparison to RLMstep.

```
> rlm1 <- rlm( HtVol ~ (Ht + BSA)^2 + Male, data = HtVol)
> summary(rlm1)
```

```
Call: rlm(formula = HtVol ~ (Ht + BSA)^2 + Male, data = HtVol)
```

```
Residuals:
```

	Min	1Q	Median	3Q	Max
	-174.3956	-39.3432	-0.1981	36.7271	305.8837

```
Coefficients:
```

	Value	Std. Error	t value
(Intercept)	178.8031	137.1983	1.3032
Ht	-1.0770	1.0845	-0.9931
BSA	-136.2376	153.4100	-0.8881
Male	28.8918	19.6876	1.4675
Ht:BSA	3.0482	0.8907	3.4221

```
Residual standard error: 56.19 on 53 degrees of freedom
```

```
> rlm2 <- rlm(HtVol~(Age+Ht+Wt+BMI+BSA+Male)
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)+I(Age*Ht)+I(Age*BMI)+I(Age*BSA)+I(Age*Ma
data=HtVol)
> summary(rlm2)
```

```
Call: rlm(formula = HtVol ~ (Age + Ht + Wt + BMI + BSA + Male) +
I(Age^2) +
```

```

I(Ht^2) + I(Wt^2) + I(BMI^2) + I(BSA^2) + I(Age * Ht) + I(Age *
BMI) + I(Age * BSA) + I(Age * Male) + I(Ht * Wt) + I(Ht *
BMI) + I(Ht * BSA) + I(Ht * Male) + I(Wt * BMI) + I(Wt *
BSA) + I(Wt * Male) + I(BMI * Male) + I(BSA * Male), data =
HtVol)
Residuals:
      Min       1Q   Median       3Q      Max
-105.208  -21.892   -2.996   26.580  131.733
.
.
.
Residual standard error: 34.96 on 33 degrees of freedom

```

Furthermore, the residual standard error for the stepAIC model is much smaller. The stepAIC model is a substantially stronger model.

2 Sensitivity Analysis

Since there is a large discrepancy between stepAIC and RLMstep, it is interesting to view how the different threshold values of RLMstep, f-to-enter and f-to-leave, effect the resulting model.

F-To-Enter	F-To-Leave	Num of Predictors	SE Residual
4	2	1	55.54
2	1	2	53.47
1	0.5	2	53.47
.5	0.25	3	52.81
.1	0.05	5	49.29
AIC Model		25	34.96

There is no reasonable f value to make RLMstep produce similar results to stepAIC

```

> RLMstep3(HtVol$HtVol,model.matrix(~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)
,data=HtVol),f.to.enter = 4,f.to.leave =2,is.bca = F)
.
.

```

```

.
$coef
      data.in.model data.in.modelHt:BSA
      53.778583      2.084142
> RLMstep3(HtVol$HtVol,model.matrix(~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)
,data=HtVol),f.to.enter = 2,f.to.leave =1,is.bca = F)
.
.
.
$coef
      data.in.model data.in.modelHt:BSA data.in.modelMale
      36.031775      2.086633      30.793939
> RLMstep3(HtVol$HtVol,model.matrix(~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)
,data=HtVol),f.to.enter = .5,f.to.leave =.25,is.bca = F)
.
.
.
$coef
      data.in.model data.in.modelHt:BSA data.in.modelMale
      36.031775      2.086633      30.793939
> RLMstep3(HtVol$HtVol,model.matrix(~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)
,data=HtVol),f.to.enter = .5,f.to.leave =.25,is.bca = F)
.
.
.
$coef
      data.in.model data.in.modelHt:BSA data.in.modelMale
      41.108304847      2.009878011      31.366775209
data.in.modelAge:Wt
      0.001215089
RLMstep3(HtVol$HtVol,model.matrix(~(Age+Ht+Wt+BMI+BSA+Male)^2
+I(Age^2)+I(Ht^2)+I(Wt^2)+I(BMI^2)+I(BSA^2)
,data=HtVol),f.to.enter = .1,f.to.leave =.05,is.bca = F)
Coefficients:
(Intercept) I(Ht * BSA) I(Age * Wt) I(Ht * Male) I(BMI * Male)
79.5992036984 2.3888971914 0.0007992141 0.6203318183 -2.6005277295
I(Ht^2)

```

-0.0054220324

```
summary(rlm.low.f)
```

.
.
.

```
Call: rlm(formula = HtVol ~ I(Ht * BSA) + I(Age * Wt) + I(Ht *  
      Male) +  
      I(BMI * Male) + I(Ht^2), data = HtVol)
```

3 Appendix: Code for RLMstep

```
RLMstep3 <- function(yinput, xinput, f.to.enter =  
  2, f.to.leave=1, bca.alpha = .95, is.bca = F){  
  if(is.bca){  
    require(boot)  
  }  
  if(bca.alpha >= 1 || bca.alpha <= 0){  
    geterrmessage("Pick bca.alpha between 0 and 1. This alpha is  
      used to construct BCA conf intervals")  
  }  
  if(f.to.enter <= f.to.leave){  
    geterrmessage("Please pick f.to.enter > f.to.leave")  
  }  
  if(xinput[,1] == 1){  
    xinput<-xinput[,-1]  
  }  
  model.list <- list()  
  max.num.pred <- ncol(xinput) # maximum number of predictors in  
    full model = # of predictors  
  n <- nrow(xinput)  
  error.final.models <- rep(NA, max.num.pred+1) # error terms for  
    each "Best" model with p predictors. +1 to include model with  
    just intercept  
  col.ind.in.data <- NULL # tracks which columns ARE in model,  
    INITIALLY NULL
```

```

col.ind.not.data <- 2:(max.num.pred+1) # tracks which columns are
    NOT in model, INITIALLY EVERYTHING OTHER THEN ITERCEPT
data.in.model <- NULL # matrix that is BUILT UP with each
    iteration as variables are ADDED to the model. INITIALLY NULL

# We initially calculate the intercept model
rlm.fit <- rlm(yinput ~ 1)
error.final.models[1] <- rss(rlm.fit)
model.list[[1]] <- rlm.fit # Include as first model
col.ind.in.data <- append(col.ind.in.data,1) # Include in the
    data in model. It is already excluded from col.not.
xinput <- cbind(1,xinput)
data.in.model <- xinput[,1]
#print(data.in.model)

#We begin the selection process
for(i in 1:max.num.pred){
  didAdd <- FALSE # bool for adding variable
  error.add.jth.var <- NULL # Max size is everything NOT in
    data null since we are starting with forward
  error.remove.jth.var <- NULL # Max size of elements we may
    remove from the model
  model.compare.rss <- error.final.models[i] # The model which we
    will use for F test comparison
  old.col.ind.data <- col.ind.in.data
  #FORWARD SELECTION
  weight1 <- NULL # Define a weight matrix
  for (j in col.ind.not.data) {
    model.build.data <- cbind(xinput[,col.ind.in.data],
      xinput[,j]) # loop through and create a dataset with last
      iteration PLUS new variable to test
    rlm.fit <- rlm(yinput ~ model.build.data-1) # Fit the model
    weight1 <- cbind(weight1,rlm.fit$w) # Gather weights
    error.add.jth.var<- append(error.add.jth.var,rss(rlm.fit)) #
      Append the errors to the vector.
  }

  ind.smallest <- which.min(error.add.jth.var)#give me the index
    of the variable which caused the biggest reduction in RSS
  min1 <- min(error.add.jth.var) # Find the min RSS

```

```

test.1 <-
  model.forward.ratio(model.compare.rss,min1,length(col.ind.in.data),weight1[,ind.s
  # F test described by Agostinelli

#print(c("Test Result Forward", test.1))
if(test.1 > f.to.enter){ # If the following condition holds we
  add the model
  var.to.add <- col.ind.not.data[ind.smallest] # Give me the
    actual variable
  col.ind.in.data <-append(col.ind.in.data,var.to.add) #Append
    new variable
  col.ind.not.data <-
    col.ind.not.data[!col.ind.not.data==var.to.add] # Remove
    from not list
  data.in.model <- xinput[,col.ind.in.data,drop=FALSE] # Input
    the variable
  model.compare.rss <- min1 # Record new comparison rss
  didAdd <- TRUE # We added
}
#Backward Selection
weight2 <- NULL #Define a second weight matrix
if(i>1 & didAdd){ #Entry Conditions
  for (j in old.col.ind.data){
    model.build.data <-
      xinput[,col.ind.in.data[col.ind.in.data!=j]] #
      Iteratively remove predictors
    rlm.fit <- rlm(yinput ~ model.build.data-1) # Regress them
    weight2 <- cbind(weight2,rlm.fit$w) # Retrieve rlm weights
    error.remove.jth.var
      <-append(error.remove.jth.var,rss(rlm.fit)) # store rss
  }
  min2 <- min(error.remove.jth.var) # Min RSS
  ind.min2 <- which.min(error.remove.jth.var) # Index min
  test.2 <-
    model.backward.ratio(min2,min1,length(old.col.ind.data),weight2[,ind.min2])
    # Agostelli test
#print(c("Test Result Backward",test.2))
if( test.2 < f.to.leave ){ # If the following holds remove

```



```

var.to.remove <- old.col.ind.data[ind.min2] # find var to
remove
col.ind.in.data <- col.ind.in.data[col.ind.in.data !=
var.to.remove] # remove it from col.in.data
col.ind.not.data <- append(col.ind.not.data, var.to.remove) #
add it to col.not
data.in.model <- xinput[, col.ind.in.data, drop=FALSE] #
update data.in.model
#Exist condition described by agostinelli to ensure
convergence
if(min2 < min1 * (1 + f.to.leave / (sum(weight2) - length(old.col.ind.data))) / (1 + f.to.ente
})
}
rlm.fit <- rlm(yinput ~ data.in.model-1) #refit model, store
it, and extract needed error term
model.list[[i+1]] <- rlm.fit ## Figure out where to store
error.final.models[i+1] <- rss(rlm.fit)
if(!didAdd){break}
else{if(i>2){if((test.1 <= f.to.enter) & (test.2 >=
f.to.leave)){break}}}
}

rlm.final.model <- rlm(yinput ~ data.in.model-1) # Compute the
final model
regression.fucntion <- function(formula, data, index){ # Boot
helper function for BCa
rlm.new <- rlm(formula, data=data[index,])
return(coef(rlm.new))
}
if(is.bca){
# Compute the BCa confidence intervals
data.frame.1 <-
cbind(data.frame(yinput), data.frame(data.in.model))
boot.rlm <- boot(data = data.frame.1, statistic
=regression.fucntion ,R = 10000, formula =yinput ~ .-1)
bca<-lapply(1:(ncol(data.in.model)), function(x) return(boot.ci(boot.out
=boot.rlm, index = x, type ="bca", conf = bca.alpha)))
}

```

```

# Store the results
df.temp.model.results <- list()
df.temp.model.results$predictor <- data.in.model
df.temp.model.results$predictor.names <- colnames(data.in.model)
if(is.bca){df.temp.model.results$bca <-
  bca}else{df.temp.model.results$bca <- "Set is.bca = TRUE"}
df.temp.model.results$coef <- rlm.final.model$coef
return(df.temp.model.results)

#df.model.results <- rbind(df.model.results,
  df.temp.model.results)
}

# Helper functions
model.forward.ratio <- function(rss.a,rss.a1,da,w){
  return((rss.a-rss.a1)/(rss.a1/(sum(w)-da-1)))
}
model.backward.ratio <- function(rss.a,rss.a1,da,w){
  return((rss.a-rss.a1)/(rss.a1/(sum(w)-da)))
}
rss <- function(rl){return(sum(rl$w*rl$resid^2) )}

```
